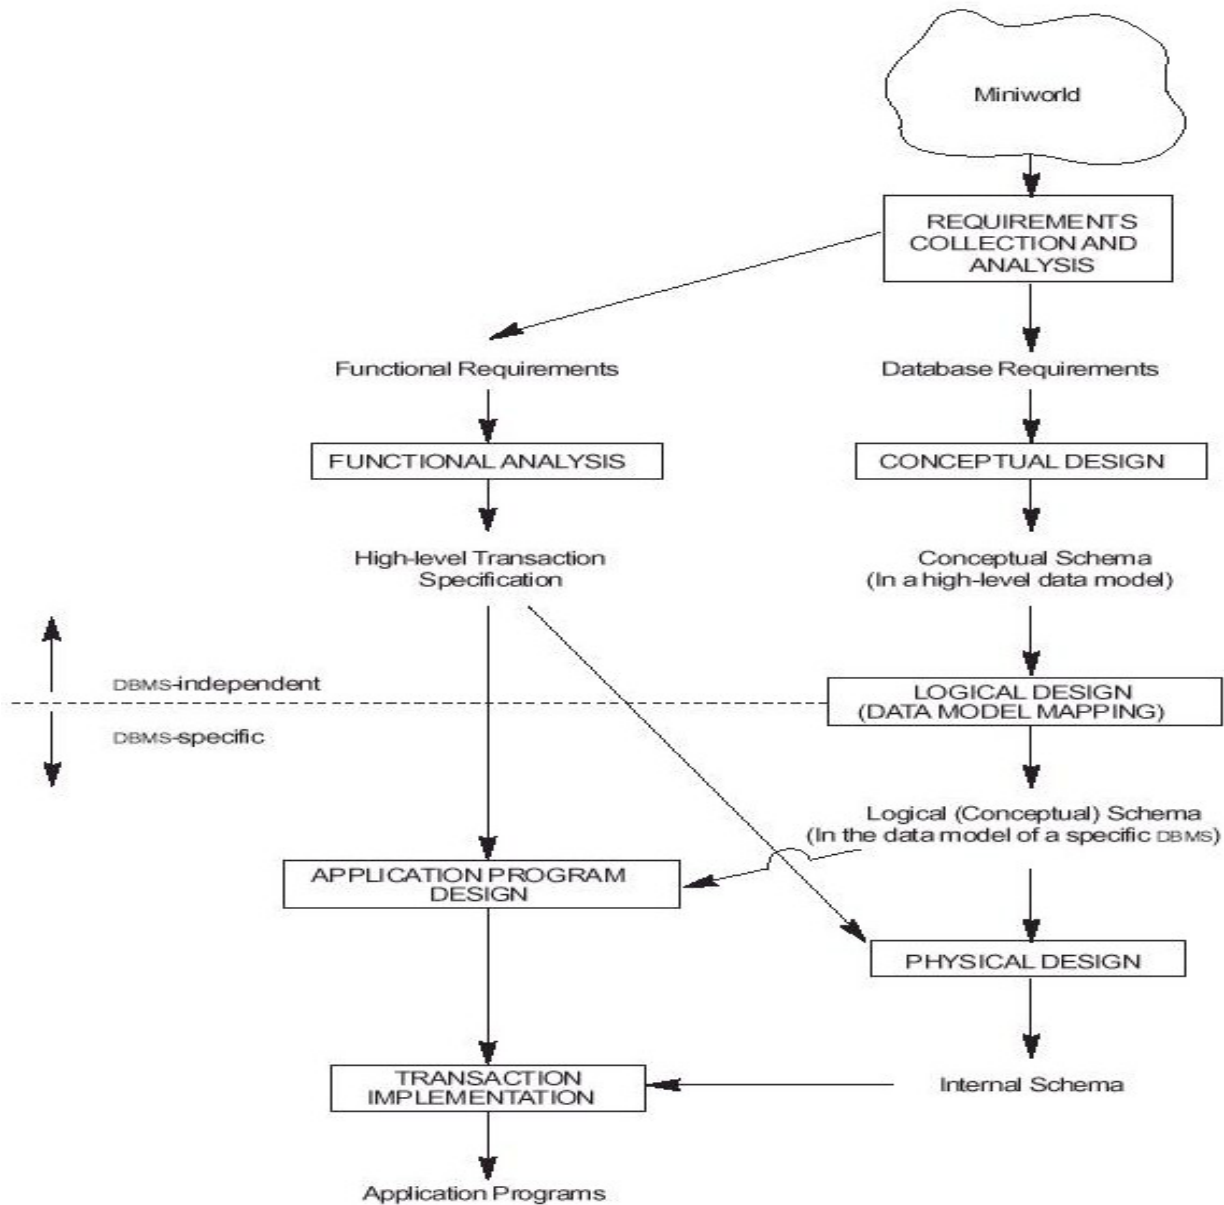


Review of E-R Data Model

Database Design

- Goal: specification of database schema
- Methodology:
 - Use E-R model to get a high-level graphical view of essential components of enterprise and how they are related
 - Convert E-R diagram to DDL
- E-R Model: enterprise viewed as set of
 - Entities
 - Relationships among entities



Overview of Database Design


- * Conceptual design: ER Model is used at this stage, which is then translated to a relational schema.
- Schema Refinement: (Normalization) Check relational schema for redundancies and related anomalies.
- Physical Database Design: Consider typical workloads and further refine (index, path access, etc) the database design.

Conceptual Design

- What are the *entities* and *relationships* in the enterprise?
- What information about these entities and relationships should we store in the database?
- What are the *integrity constraints* that hold?
- Represent this information pictorially in *ER diagrams*, then map ER diagram into a relational schema.

ER Diagrams

- **Entities**: real world “objects”, = “thing.”
 - e.g. Students, Courses, Routes, Climbers. Entity sets are drawn as rectangles.
- Entity Type: set of “similar” entities/objects.
- **Attributes**: properties typically used as column names, similar to fields of a struct
 - e.g. Name, Age, Height. Drawn as ovals.




Relationships

- Connect two or more entity sets.
- Represented by diamonds.




Drawing Entity Type

- The term “entity” and “entity type” are often confused. Boxes represent *sets* of entities.
- To draw an entity set we connect it with its attributes and indicate the key by underlining it:



Drawing Relationships

- Relationship types are represented by diamonds.
- They connect the entities they relate, and may additionally have attributes.



Structural Constraints – one way to express semantics of relationships




Structural constraints on relationships:

- **Cardinality ratio** (of a binary relationship): 1:1, 1:N, N:1, or M:N
- **Participation constraint** (on each participating entity type)

Structural Constraints – one way to express semantics of relationships

Structural constraints on relationships:

- **Cardinality ratio** (of a binary relationship): 1:1, 1:N, N:1, or M:N



- **Participation constraint** (on each participating entity type)

Structural Constraints – one way to express semantics of relationships

Structural constraints on relationships:

- **Cardinality ratio** (of a binary relationship): 1:1, 1:N, N:1, or M:N

SHOWN BY PLACING APPROPRIATE NUMBER ON THE LINK.

- **Participation constraint** (on each participating entity type):
 - total (called *existence dependency*)
 - partial.

SHOWN BY DOUBLE LINING THE LINK

Alternative (min, max) notation for structural constraint

- Specified on *each participation* of an entity type E in a relationship type R
- Specifies that each entity e in E participates in *at least* min and *at most* max relationship instances in R
- Default(no constraint): min=0, max=n
- Must have $\text{min} \leq \text{max}$, $\text{min} \geq 0$, $\text{max} \geq 1$
- Derived from the knowledge of mini-world constraints

Alternative (min, max) notation for structural constraints:

Examples:

- A degree program can have any # of student majoring in and a student can major in at most 4 different programs at the same school.

Alternative (min, max) notation for structural constraints:

Examples:

- A degree program can have any # of student majoring in and a student can major in at most 4 different programs at the same school.
 - Specify (0,4) for participation of Student in MajorsIn
 - Specify (0,N) for participation of Program in MajorsIn.

Alternative (min, max) notation for structural constraints:

Examples:

- A degree program can have any # of student majoring in and a student can major in at most 4 different programs at the same school.
 - Specify (0,4) for participation of Student in MajorsIn
 - Specify (0,N) for participation of Program in MajorsIn.
- A program can be offered to any number of student, but a student majors in exactly one program.

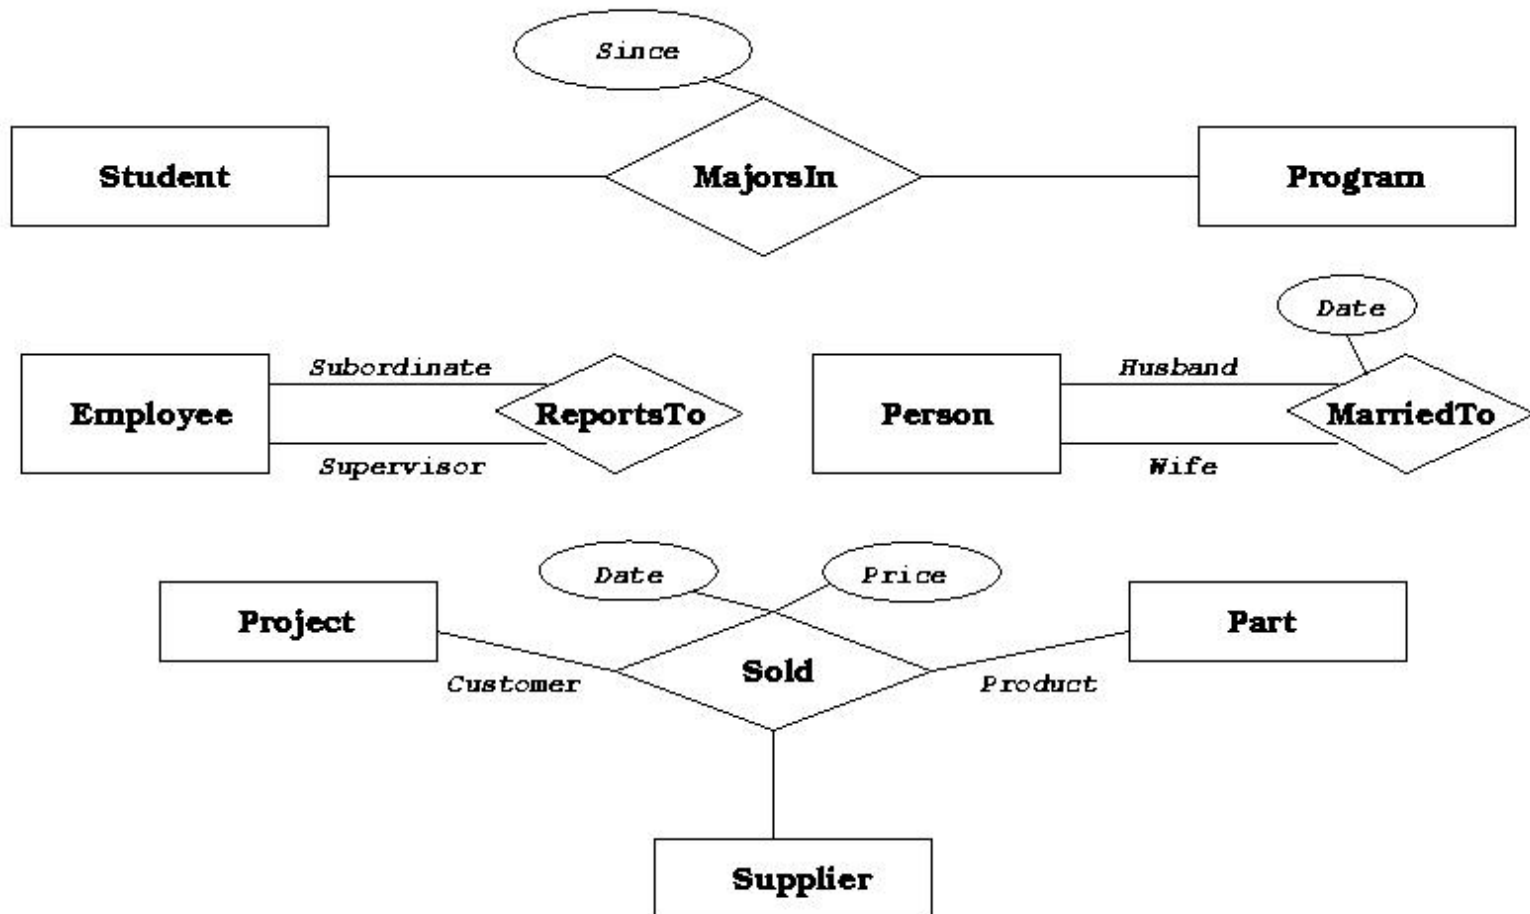
Alternative (min, max) notation for structural constraints:

Examples:

- A degree program can have any # of student majoring in and a student can major in at most 4 different programs at the same school.
 - Specify (0,4) for participation of Student in MajorsIn
 - Specify (0,N) for participation of Program in MajorsIn.
- A program can be offered to any number of student, but a student majors in exactly one program.
 - Specify (1,1) for participation of Student in MajorsIn.
 - Specify (0,N) for participation of Program in MajorsIn.

Graphical Representation

- Roles are edges labeled with role names (omitted if role name = name of entity set). Most attributes have been omitted.

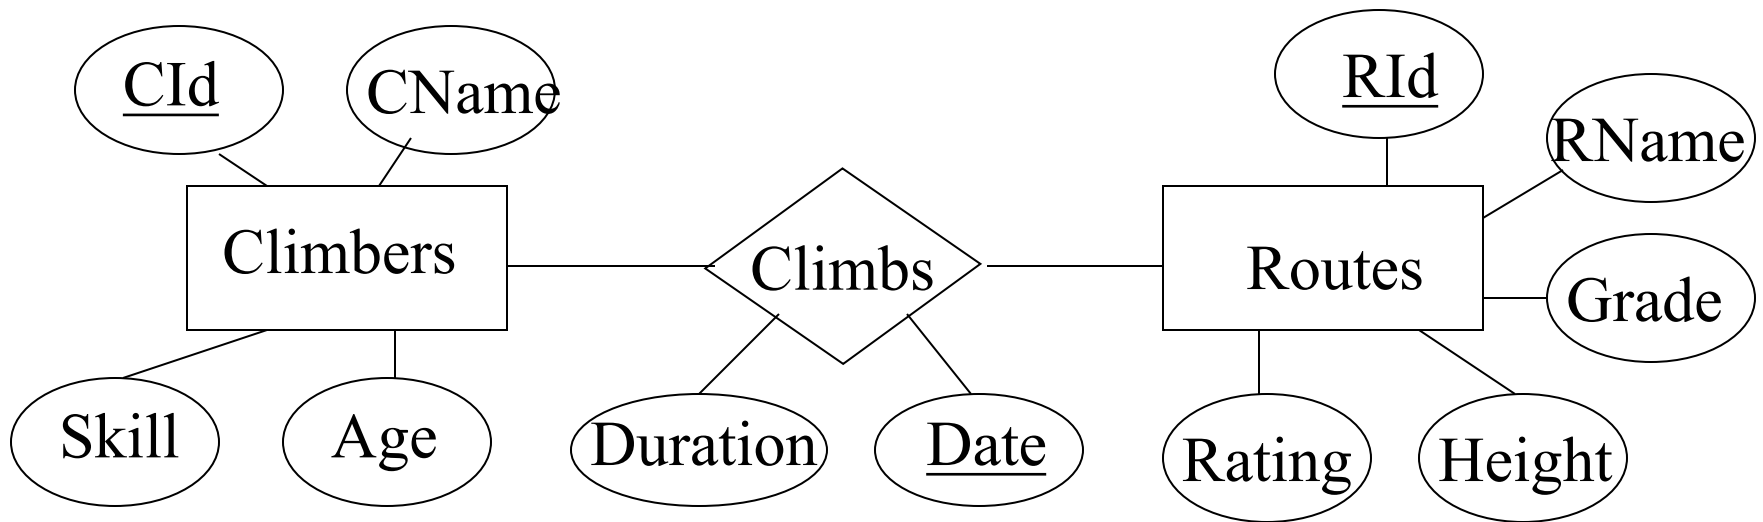


The whole diagram

Note:

lines connect entities to attributes and relationships to entities and to attributes.

No connection between attributes to attributes, entities to entities or relationships to relationships.



Weak Entity Types


Sometimes an E.T. E 's key comes not (completely) from its own attributes, but from the keys of one or more E.T.'s to which E is linked by a *supporting (identifying)* many-one relationship.

- Called a *weak* E.S.
- Represented by putting double rectangle around E and a double diamond around each supporting relationship.
- Many-one-ness of supporting relationship (includes 1-1) essential.
 - With many-many, we wouldn't know which entity provided the key value.
- “Exactly one” also essential, or else we might not be able to extract key attributes by following the supporting relationship.

Weak Entity Types

- An entity that does not have a key attribute
- A weak entity must participate in an identifying relationship type with an owner or identifying entity type
- Entities are identified by the combination of:
 - A partial key of the weak entity type
 - The particular entity they are related to in the identifying entity type


Example: Logins (Email Addresses)



Example: Logins (Email Addresses)


Login name = user name + host name, e.g., `abc@cs.umn.edu`.

- A “login” entity corresponds to a user name on a particular host, but the passwd table doesn’t record the host, just the user name, e.g., `abc`.
- Key for a login = the user name at the host (which is unique for that host only) + the IP address of the host (which is unique globally).




Example: Chain of “Weakness”

Consider IP addresses consisting of a primary domain (e.g., edu), subdomain (e.g., umn), and host (e.g., cs).



Example: Chain of “Weakness”


Consider IP addresses consisting of a primary domain (e.g., edu), subdomain (e.g., umn), and host (e.g., cs).



- Key for primary domain = its name.

Example: Chain of “Weakness”


Consider IP addresses consisting of a primary domain (e.g., edu), subdomain (e.g., umn), and host (e.g., cs).



- Key for primary domain = its name.
- Key for secondary domain = its name + name of primary domain.


Example: Chain of “Weakness”

Consider IP addresses consisting of a primary domain (e.g., edu), subdomain (e.g., umn), and host (e.g., cs).

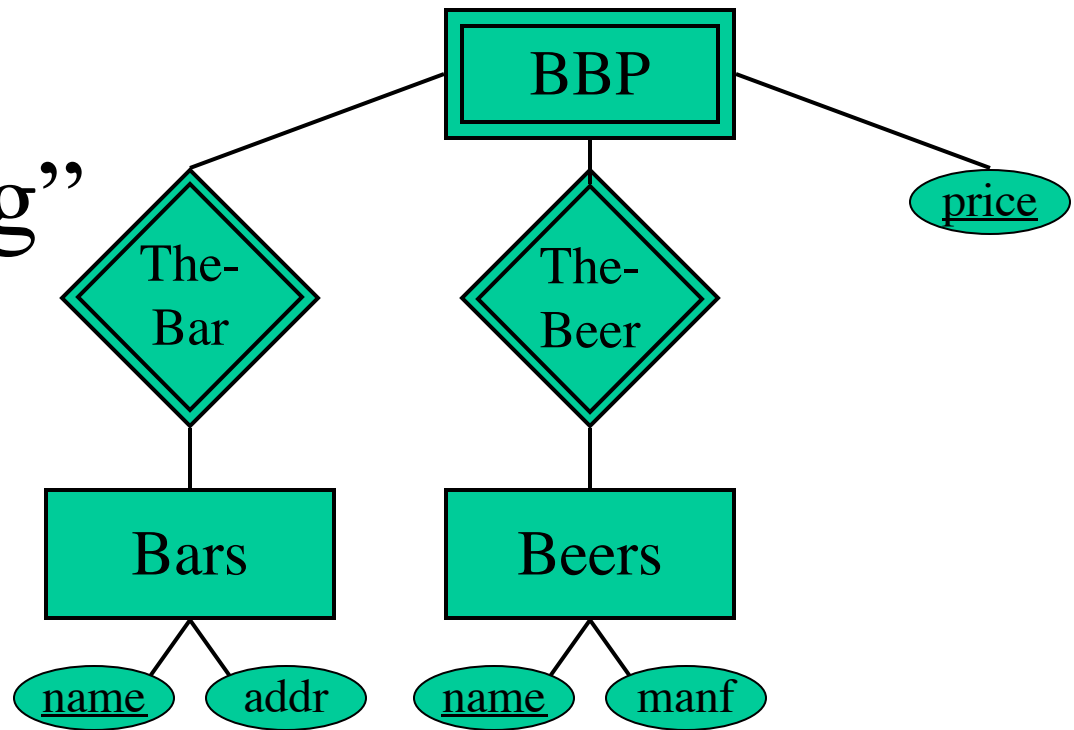


- Key for primary domain = its name.
- Key for secondary domain = its name + name of primary domain.
- Key for host = its name + key of secondary domain = its name + name of secondary domain + name of primary domain.

All “Connecting”
Entity Sets
Are Weak



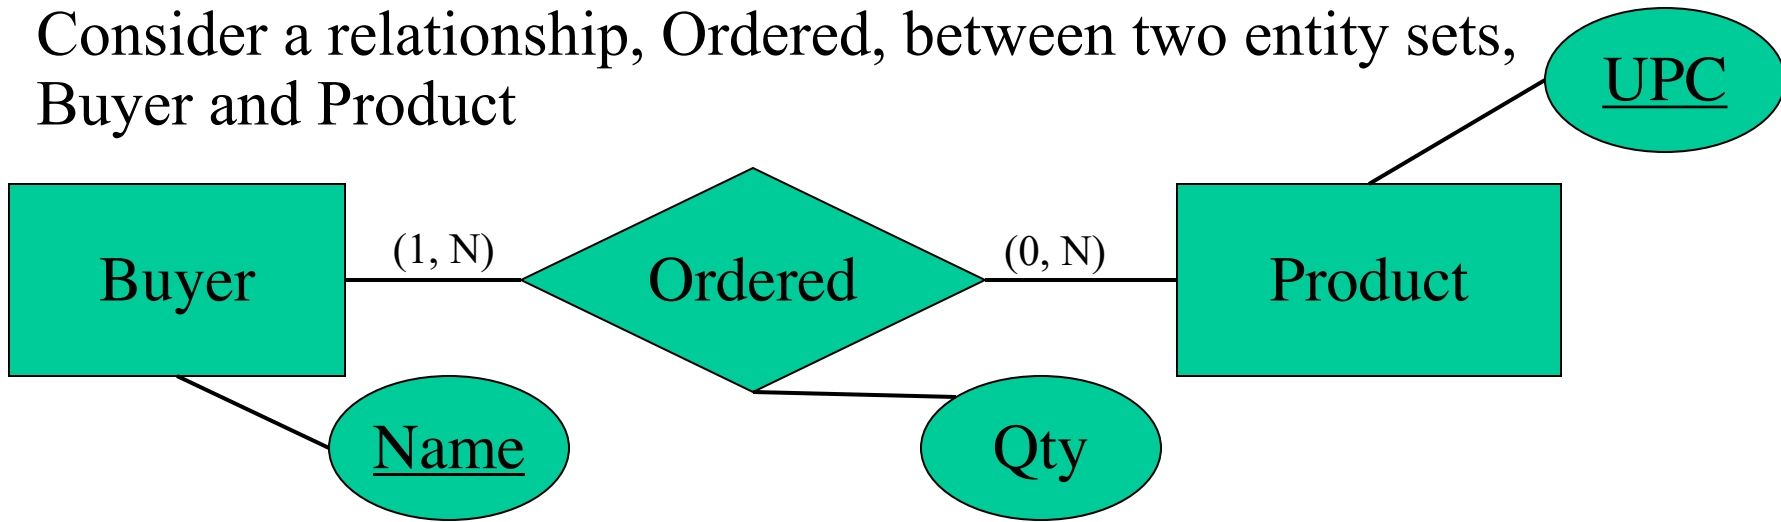
All “Connecting”
Entity Sets
Are Weak



- In this special case, where bar and beer determine a price, we can omit **price** from the key, and remove the double diamond from **ThePrice**.
- Better: **price** is attribute of **BBP**.

Relationship Vs Weak Entities

- Consider a relationship, Ordered, between two entity sets, Buyer and Product



- How can we add Shipments to track Buyer's order status?

Design Issues

Setting: client has (possibly vague) idea of what they want.
You must design a database that represents these thoughts and only these thoughts.

Avoid redundancy

= saying the same thing more than once.

- Wastes space and encourages inconsistency.

Example

Design Issues

Setting: client has (possibly vague) idea of what they want.
You must design a database that represents these thoughts
and only these thoughts.


Avoid redundancy

= saying the same thing more than once.

- Wastes space and encourages inconsistency.


Example

Good:




Example

Bad: repeats manufacturer address for each beer they manufacture.



Bad: manufacturer's name said twice.



Use Schema to Enforce Constraints

- The design *schema* should enforce as many constraints as possible.
 - Don't rely on future data to follow assumptions.

Example


- If registrar wants to associate only one instructor with a course,
- don't allow sets of instructors and count on departments to enter only one instructor per course.

Entity Sets Vs. Attributes


You may be unsure which concepts are worthy of being entity sets, and which are handled more simply as attributes.

- Don't create needless entity sets to make project “larger.”

Bad:



Good:




Intuitive Rule for E.S. Vs. Attribute

Make an entity set only if it either:

1. Is more than a name of something; *i.e.*, it has nonkey attributes or relationships with a number of different entity sets, or
2. Is the “many” in a many-one relationship.

Example

The following design illustrates both points:



- *Manfs* deserves to be an E.S. because we record *addr*, a nonkey attribute.
- *Beers* deserves to be an E.S. because it is at the “many” end.
 - If not, we would have to make “set of beers” an attribute of *Manfs* – something we avoid doing, although some may tell you it is OK in E/R model.

Don't Overuse Weak E.T.

- There is a tendency to feel that no E.T. has its entities uniquely determined without following some relationships.
- However, in practice, we almost always create unique ID's to compensate: social-security numbers, VIN's, etc.

LIMITATIONS OF THE ER MODEL

- No relationship may be defined between an entity type and a relationship type
- No relationship may be defined between an entity type and a collection of entity types from which any one type may participate
 - e.g. Entity type1 : POLICY-HOLDER may be an individual, multiple individuals, one organization, or many organizations
 - Entity type2 : POLICY

LIMITATIONS OF THE ER MODEL

- No constraints (e.g., exclusion, co-existence etc.) among relationship types
 - NIAM model, UML class diagrams allow them
 - UML (Unified Modeling Language) is a standard way to specify, construct, and document systems that use object-oriented code such as Java, C++

Extended Entity-Relationship (EER)Model

- Incorporates Set-subset Relationships
- Incorporates Generalization Hierarchies
- Constraints:
 - Coverage Constraints: partial vs. total
 - Disjointedness Constraint: disjoint vs. overlapping

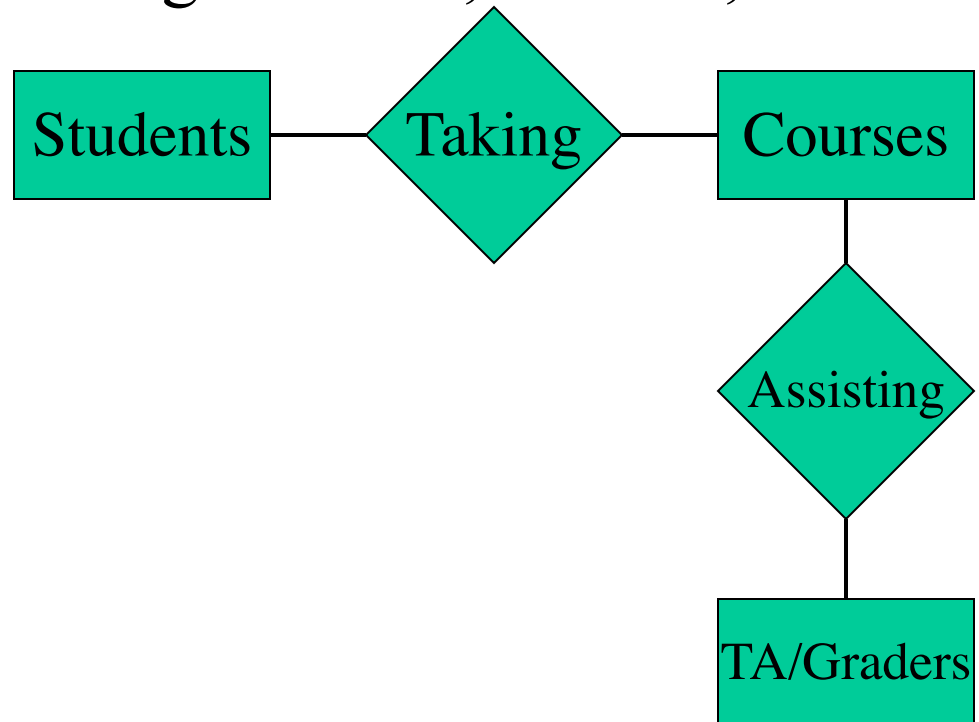
Rest of the lecture discusses HOW THE ER
MODEL CAN BE EXTENDED WITH

- **Set-subset** relationships
- **Generalization** Hierarchies
- how we can impose **further notation** on them

Multi-way Relationships

Usually binary relationships (connecting two E.T.) suffice.

- However, there are some cases where three or more E.T. must be connected by one relationship.
- Example: relationship among students, courses, TA's (and graders).




Example

- This binary relationship OK if each TA (or grader) is a TA of all students. Connection student-TA is *only* via the course.
- But what if students were divided into sections, each headed by a TA?




Example

- This binary relationship OK if each TA (or grader) is a TA of all students. Connection student-TA is *only* via the course.
- But what if students were divided into sections, each headed by a TA?
 - Then, a student in CS 6360 would be related to only one of the TA's for CS6360. Which one?
- Need a 3-way relationship to tell.




Students	Courses	TAs
Ann	CS6360	Jan
Sue	CS6360	Pat
Bob	CS6360	Jan
Linda	CS6360	Mike
...

TERNARY RELATIONSHIPS



TERNARY VS. BINARY RELATIONSHIPS




TERNARY RELATIONSHIP- Instance Diagram


SUPPLIER

SUPPLY

PROJECT

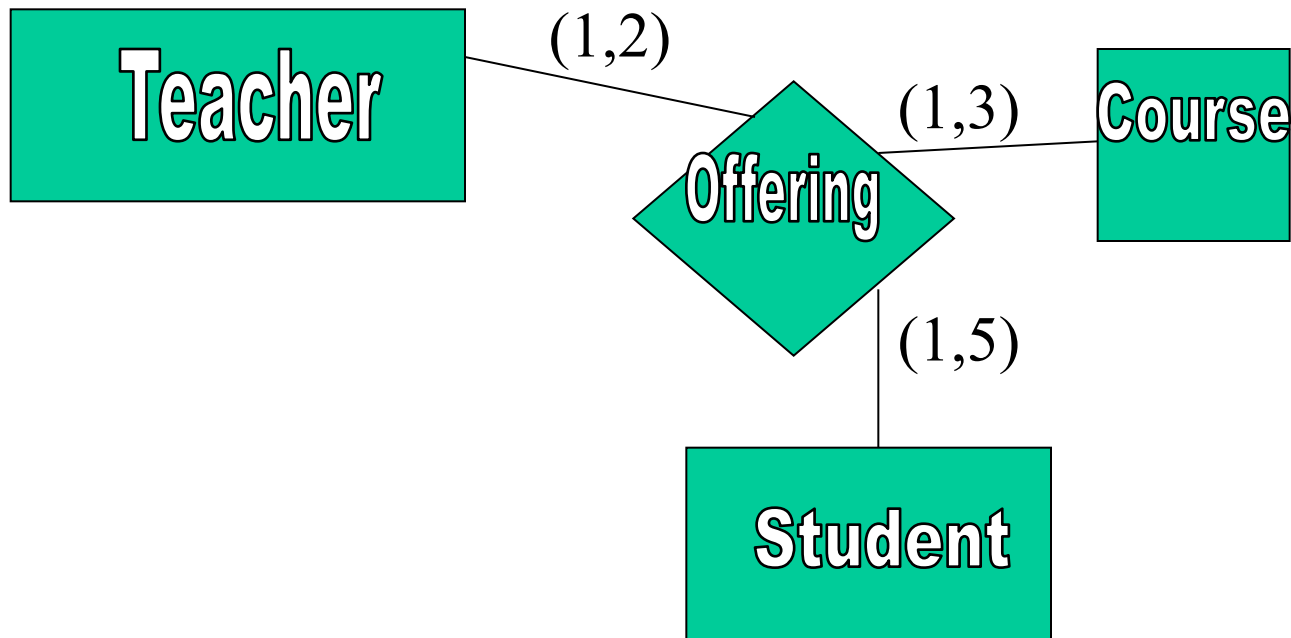


Problem with constraints on higher order relationship types



What does it mean to put $m:n:p$ on the three arms of the relationship ? It is essentially meaningless.

The (min,max) notation for higher order relationship type constraints



A Teacher can offer min 1 and max 2 Offerings


A Course may have 1 to 3 Offerings

A Student may enroll in from 1 to 5 Offerings

RECURSIVE RELATIONSHIP SUPERVISION

EMPLOYEE

WORKS_FOR



Some of the Currently Available Automated Database Design Tools

COMPANY	TOOL	FUNCTIONALITY
Embarcadero Technologies	ER Studio	Database Modeling in ER and IDEF1X
	DB Artisan	Database administration and space and security management
Oracle	Developer 2000 and Designer 2000	Database modeling, application development
Popkin Software	System Architect 2001	Data modeling, object modeling, process modeling, structured analysis/design
Platinum Technology	Platinum Enterprise Modeling Suite: Erwin, BPWin, Paradigm Plus	Data, process, and business component modeling
Persistence Inc.	Pwertier	Mapping from O-O to relational model
Rational	Rational Rose	Modeling in UML and application generation in C++ and JAVA
Rogue Ware	RW Metro	Mapping from O-O to relational model
Resolution Ltd.	Xcase	Conceptual modeling up to code maintenance
Sybase	Enterprise Application Suite	Data modeling, business logic modeling
Visio	Visio Enterprise	Data modeling, design and reengineering Visual Basic and Visual C++

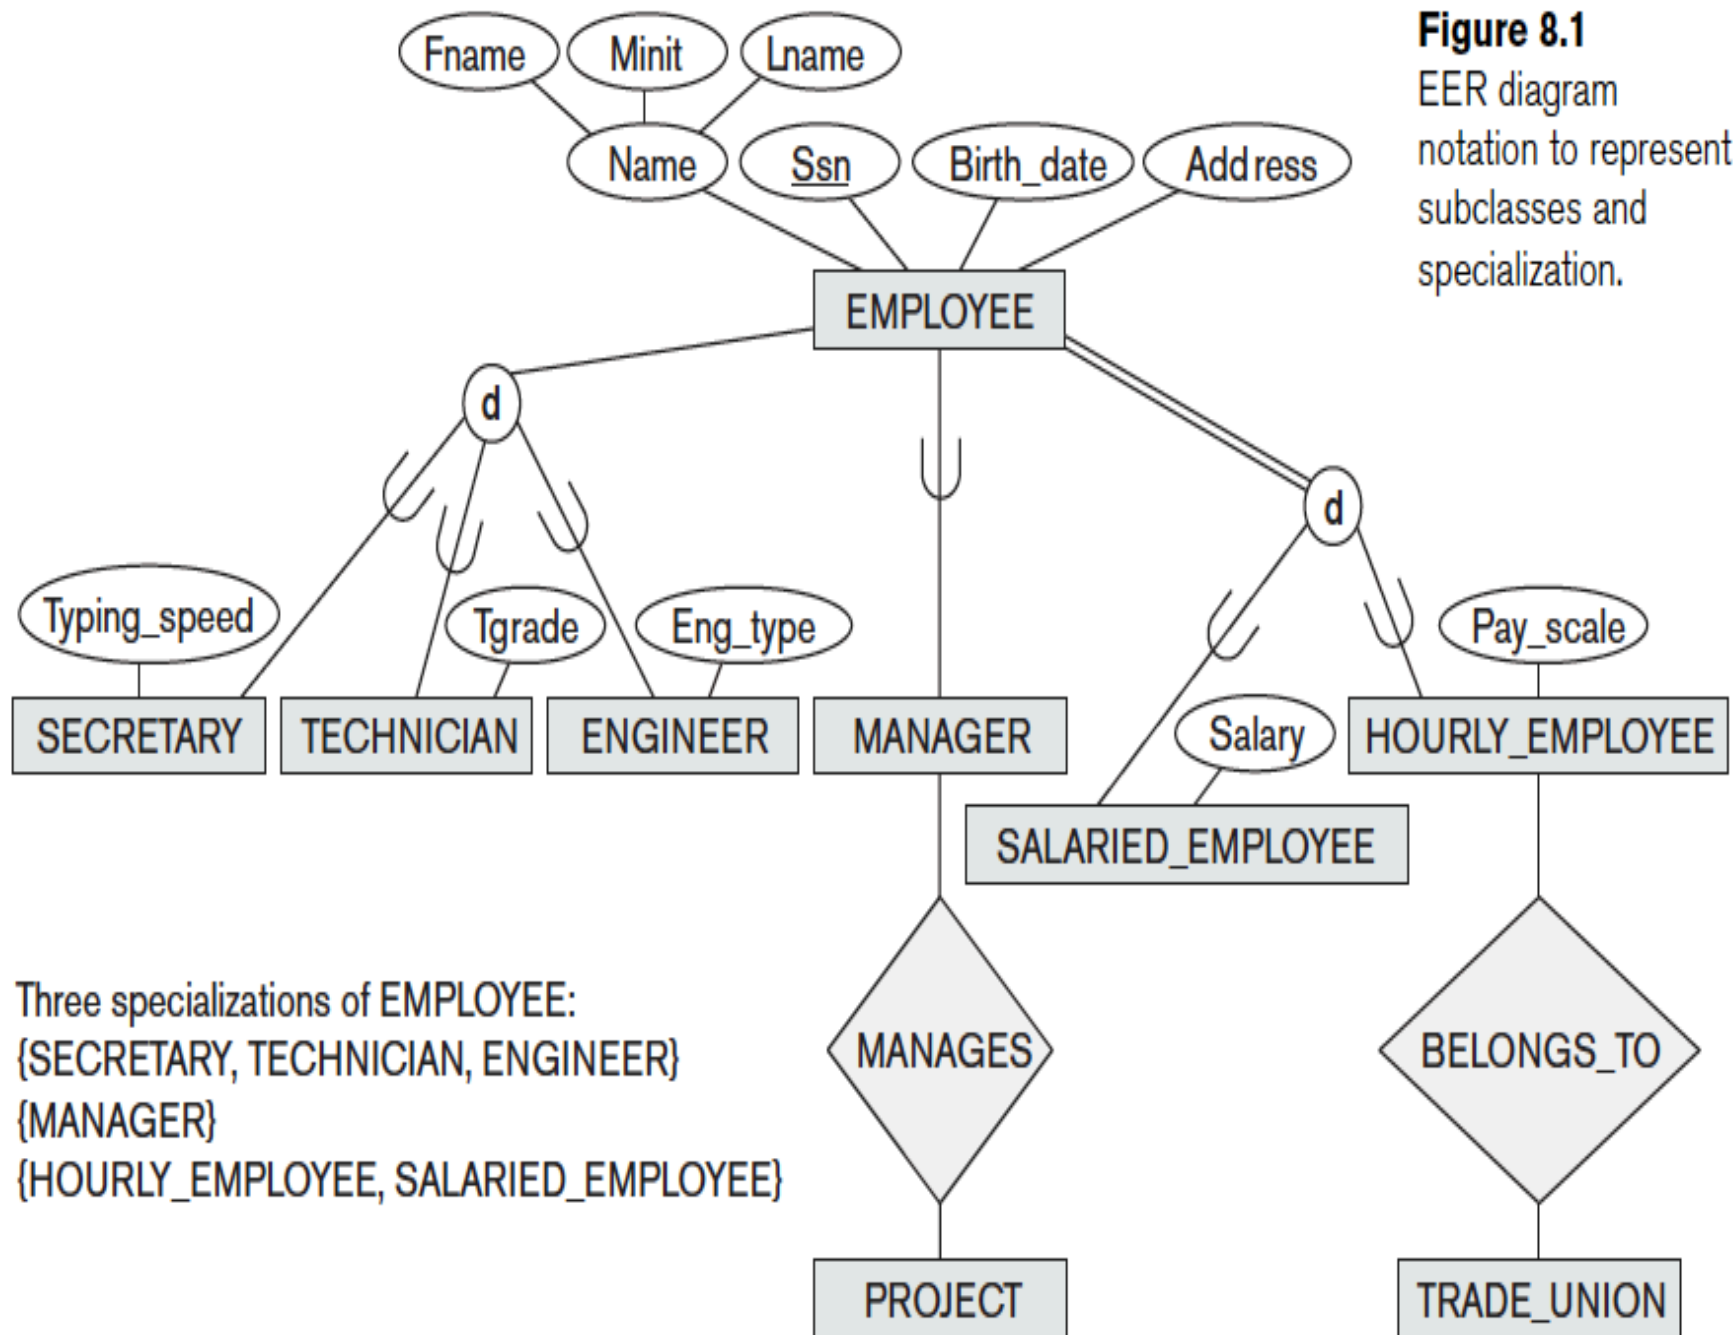


Figure 8.1
EER diagram notation to represent subclasses and specialization.

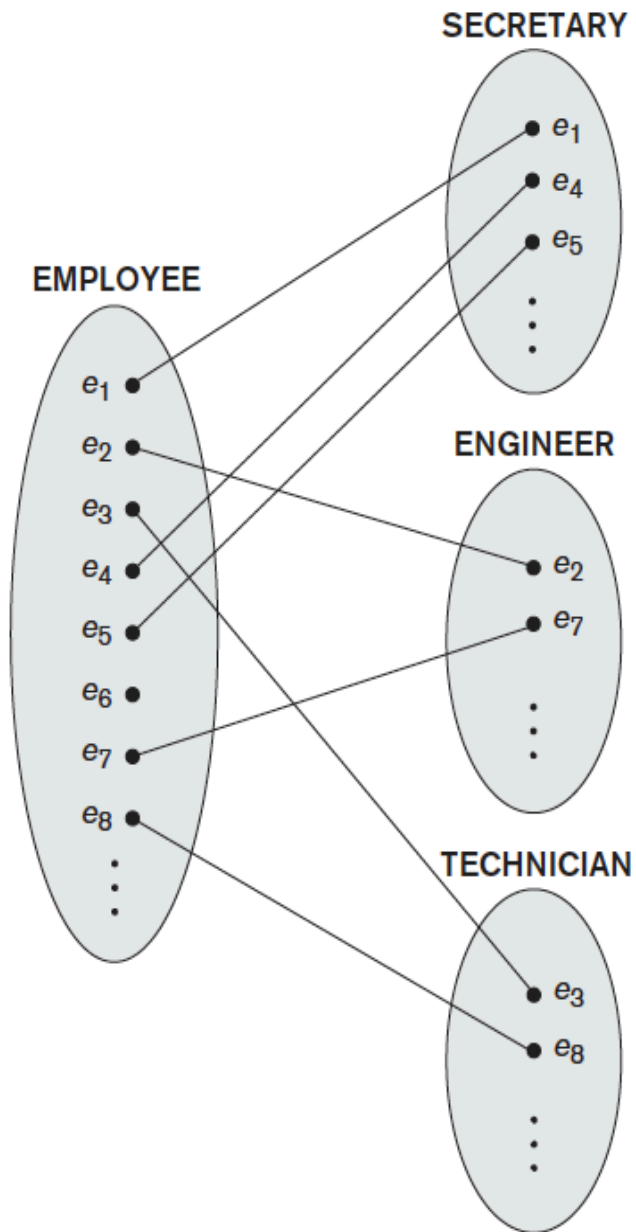


Figure 8.2
Instances of a specialization.

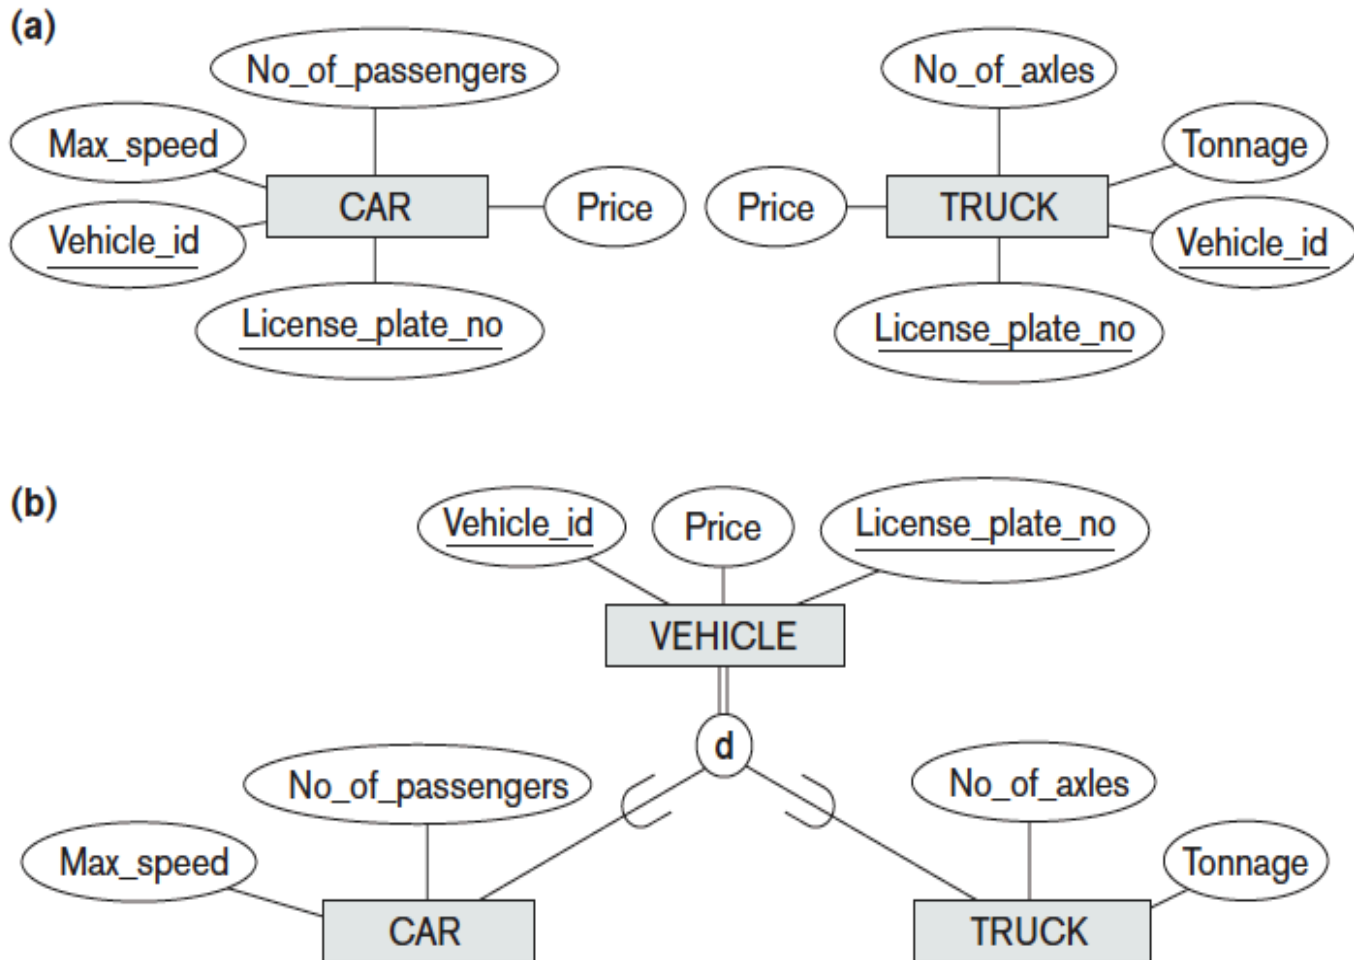
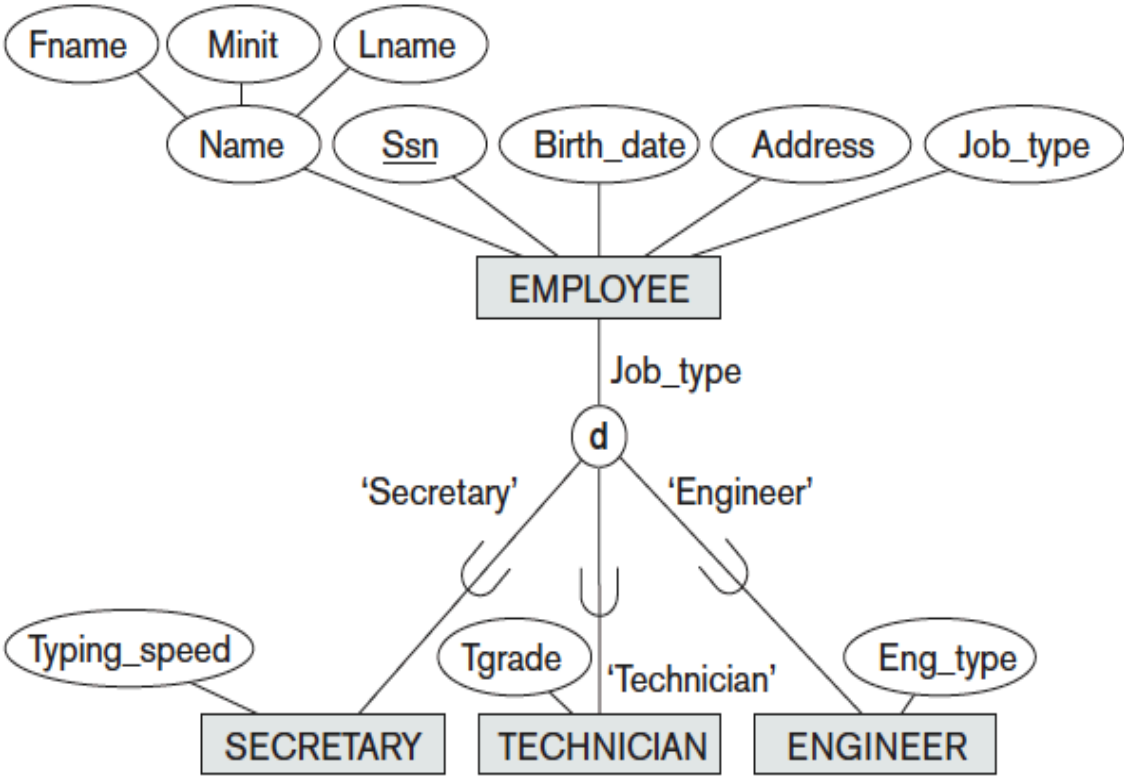


Figure 8.3

Generalization. (a) Two entity types, CAR and TRUCK. (b) Generalizing CAR and TRUCK into the superclass VEHICLE.

Figure 8.4

EER diagram notation
for an attribute-defined
specialization on
Job_type.



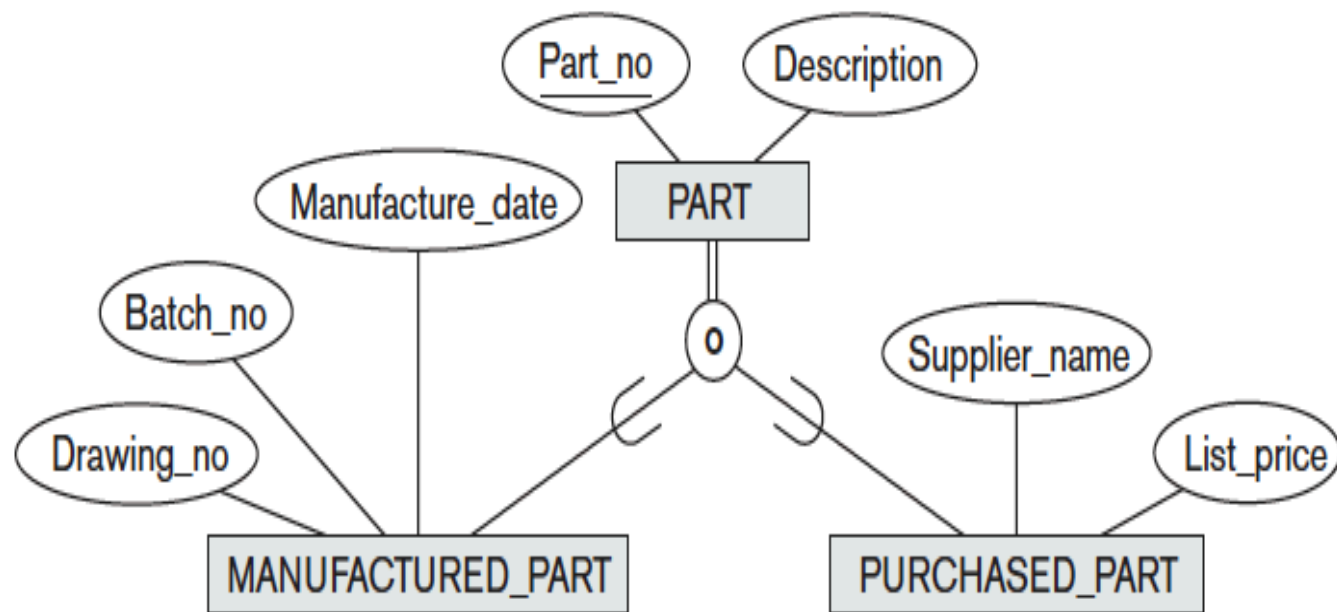
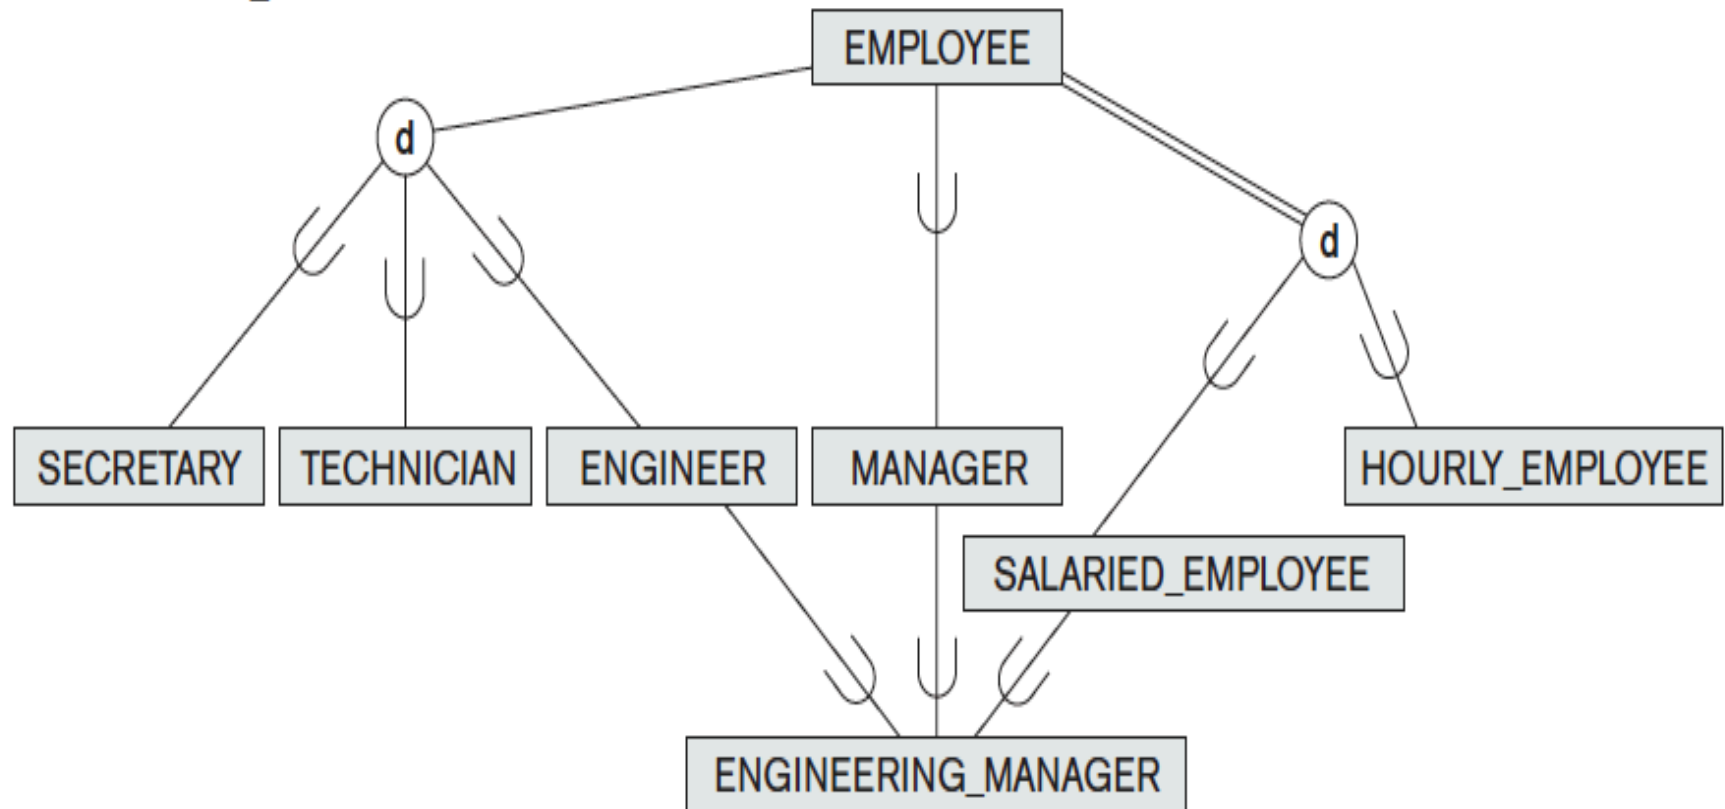


Figure 8.5

EER diagram notation
for an overlapping
(nondisjoint)
specialization.

Figure 8.6

A specialization lattice with shared subclass
ENGINEERING_MANAGER.



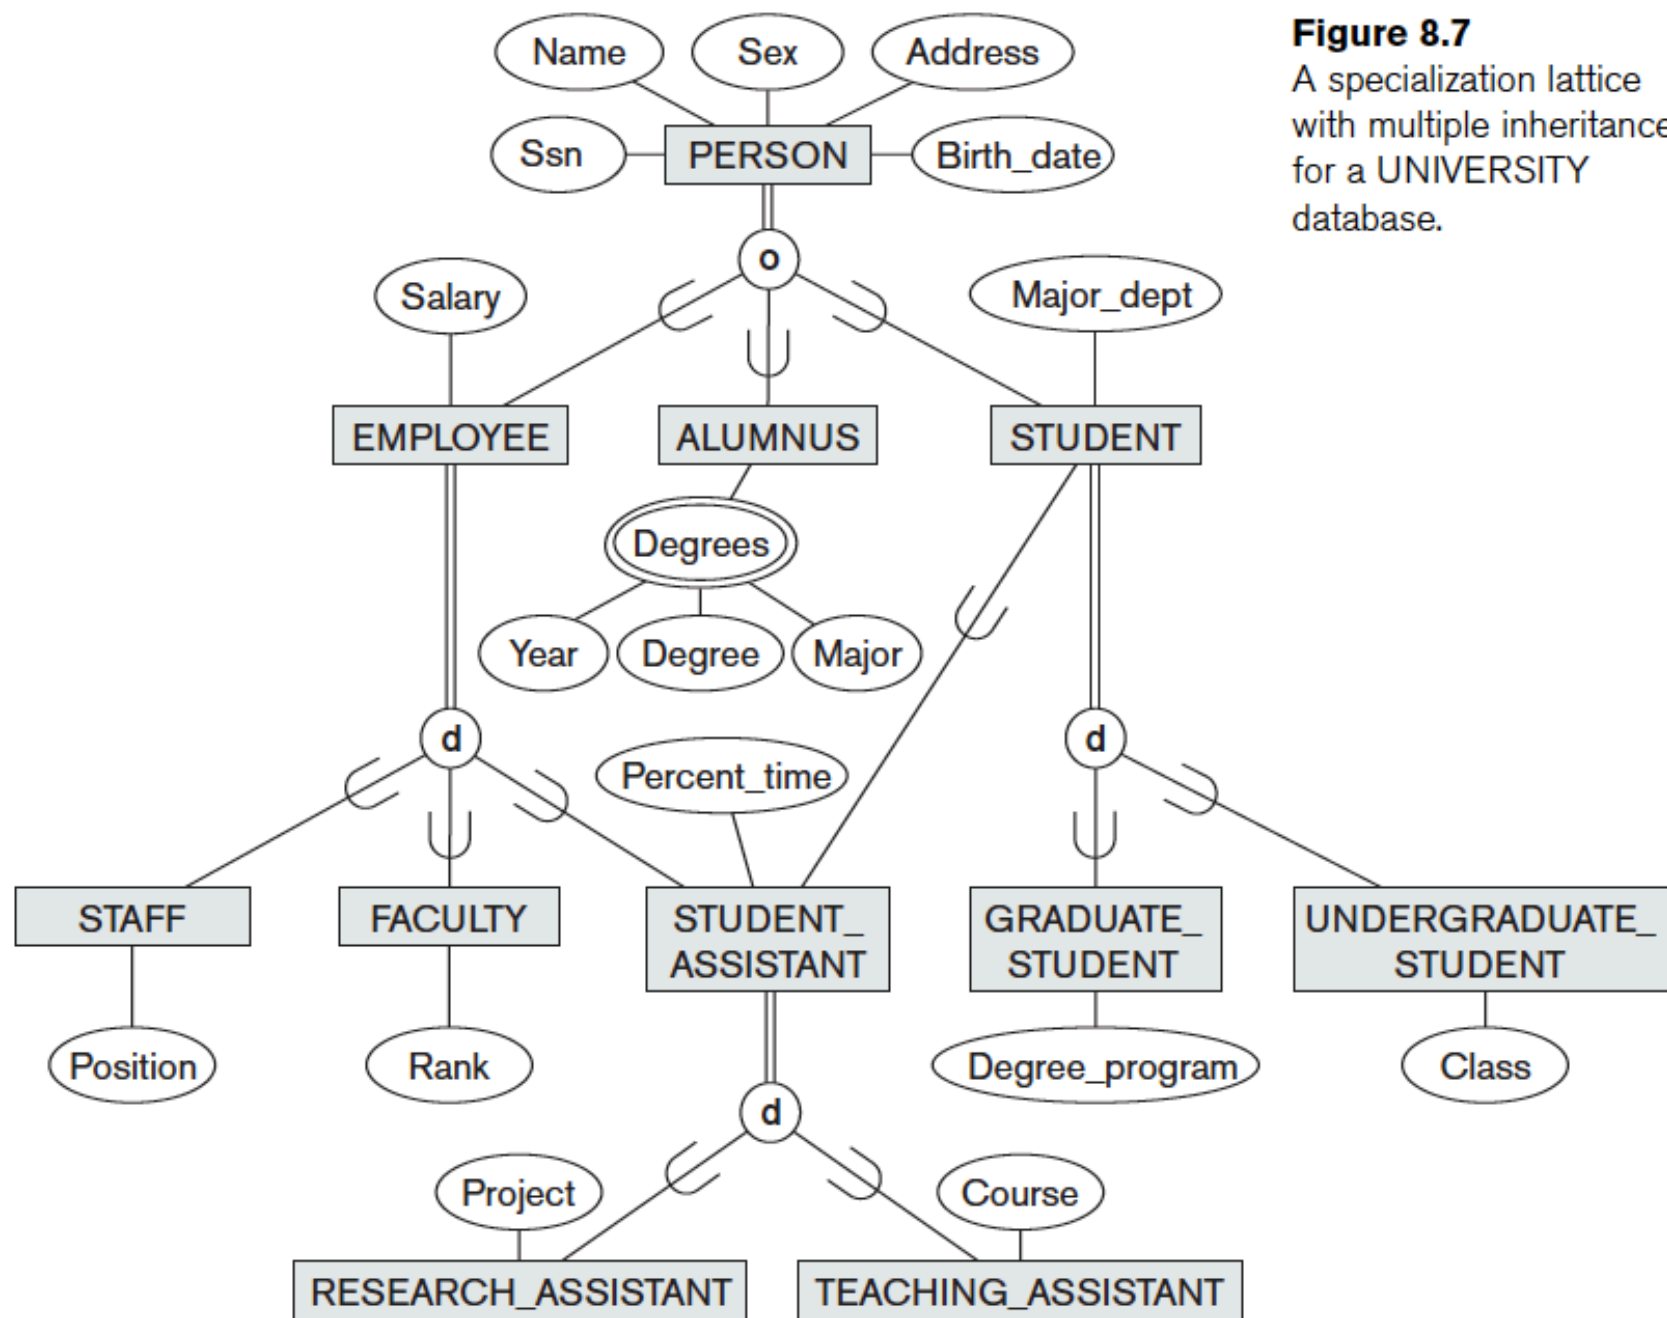


Figure 8.7

A specialization lattice with multiple inheritance for a UNIVERSITY database.

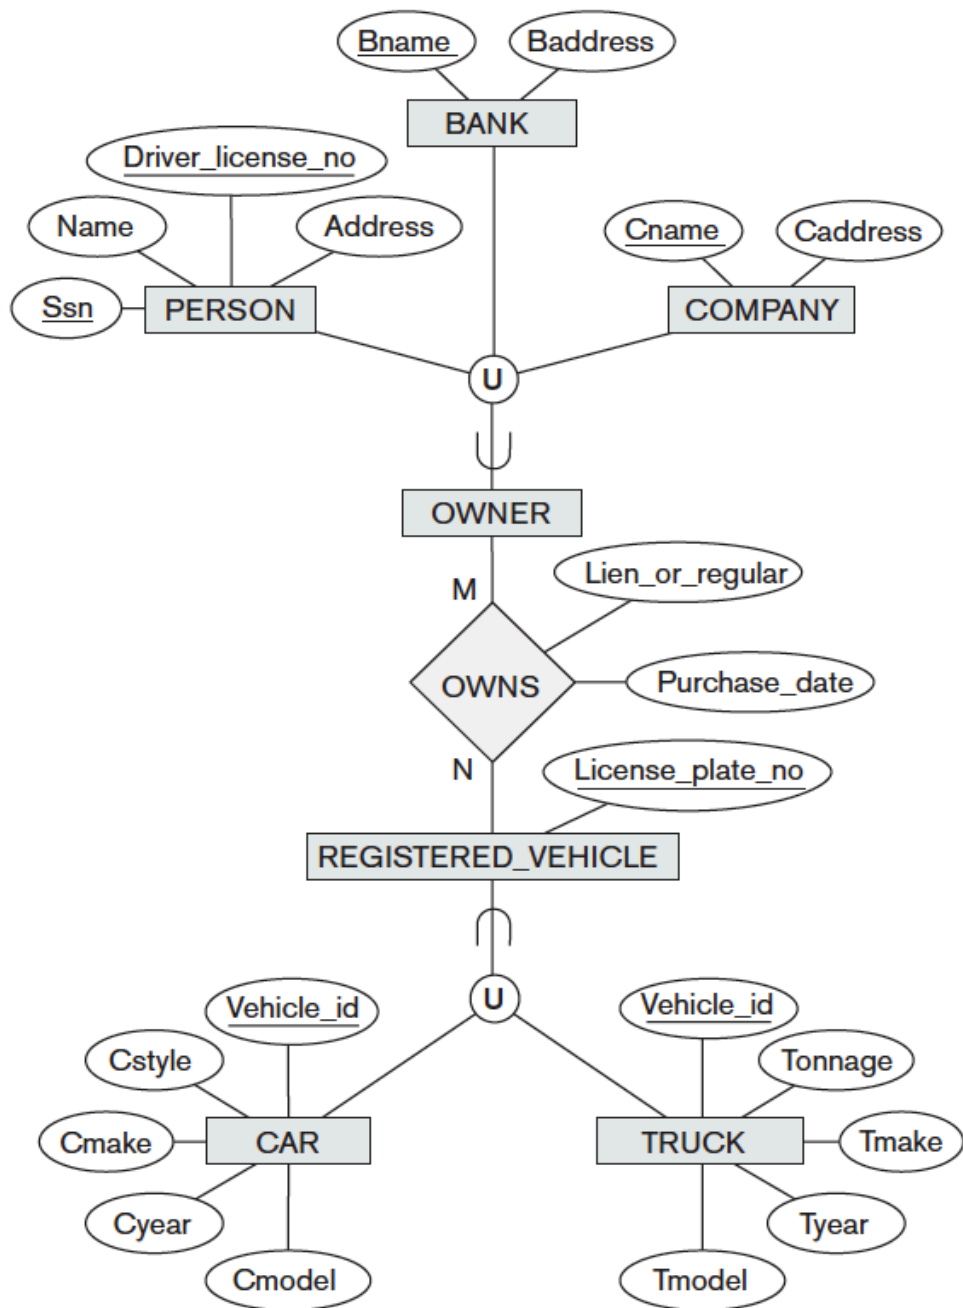


Figure 8.8
Two categories (union types): OWNER and REGISTERED_VEHICLE.