

Concurrent Control Techniques

Chapter 22

Review:

ACID Properties of Transactions

- Transaction execution must maintain the correctness of the database model
- Therefore additional requirements are placed on the execution of transactions beyond those placed on ordinary programs
 - **A**tomicity
 - **C**onsistency
 - **I**solation
 - **D**urability

Serializable Schedules

- The concurrent schedule $S: r_1(x) w_2(z) w_1(y)$ is equivalent to the serial schedules of $T1$ and $T2$ in either order:
 - $T1, T2: r_1(x) w_1(y) w_2(z)$ and
 - $T2, T1: w_2(z) r_1(x) w_1(y)$

since operations of distinct transactions on different data items commute. Hence, S is a serializable schedule

Classification of CCTs

- 22.1 2PL
 - lock types, basic/2PL/strict
 - deadlocks
- 22.2 Timestamp Ordering
 - timestamps, ordering algo
- 22.3 Multiversion CCT
 - w time-stamps, 2-phase
- 22.4 Validation (Optimistic) based CCT

What is concurrent control?

- Def.: Protocols, Algorithms, Techniques to avoid problems
 - resulting from concurrent operations on shared data
 - Problems of lost update, incorrect summary ...
- A Well-known problem in O.S., Client-Server, Shared Resource
 - Ex. Concurrent processes/threads in Operating Systems
 - Mutual Exclusion Protocols: semaphores, critical sections

Common Concurrency Control Techniques (CCTs)

- Goal:
 - Quickly ensure serializability
 - Fast filter to avoid violation of serializability
- may rule many possible serializable schedule
- Prescribes a Protocol to be observed by each transaction
- $\langle \text{current state, new op. request} \rangle \rightarrow \text{grant/delay op, abort T}$

Simple CCTs

- 2 phase locking (2PL) protocol for transactions
 - Schedule: order T_i by the order of lock acquisition
- TSO: item-based violation test at read/write
 - Schedule: order T_i by their time-stamps
- *Other schemes*
 - Optimistic, 3 phases (read, validate, write)
 - Validate phase = conflict? w/committed/validated T_j
 - Multiversion - uses version semantics

Lock Based CCT : Basic Definitions

- *LOCK[X] = a variable describing status of data item X*
 - wrt possible applicable operations without conflict.
- *Lock(X, op) = permission to Ti to perform op on X*
 - NOTE Convention: LOCK(X) = variable, Lock(X) = permission
- *Operations on LOCK(X) requested by Ti to Lock Manager*
 - lock_item(X), lock_item(X, read), unlock(X), ...
 - Protocol: Ti needs Lock(X, op) to perform op(X)
 - Lock Manager: (operation on LOCK(X)) $\rightarrow \{ \text{grant, postpone, deny} \}$
 - See Fig. 22.1, 22.2 (pp 779-781) for algorithms

Types of Locks

- binary locks:
 - Binary values (states): unlock= 0, lock=1
 - ops.: lock_item(X), unlock_item(X)
 - See protocol rules on pp 779
- Read/Write (Shared/Exclusive) locks:
 - 3 Values:
 - unlocked, read_locked (shared_locked), write_locked(exclusive_locked).
 - 3 ops:
 - read_lock(X) is shared: Multiple read locks can be allowed on X.
 - write_lock(X) exclusive, no other concurrent lock on X
 - unlock(X)
 - See pp779-782 the list of locking rules
- Conversion of Locks
 - Relax rule 4 and 5
 - Can convert the lock from one lock state to another
 - Upgrade lock: read_lock(X) \rightarrow write_lock(X)
 - Downgrade: write_lock(X) \rightarrow read_lock(X)

Implementing Serializability: Two-Phase Locking


- Locks are associated with each data item
- A transaction must acquire a read (shared) or write (exclusive) lock on an item in order to read or write it
- A write lock on an item *conflicts* with all other locks on the item; a read lock conflicts with a write lock
- If $T1$ requests a lock on x and $T2$ holds a conflicting lock on x , $T1$ must wait

Lock based CCT: 2 PL protocol


- Two-phase Locking (2PL): T_i follows 2PL if
 - all locking ops precede the first unlock op in T
- Two Phases relate to the text/execution of T
 - Phase 1: Expand / Grow the set of permissions (Locks)
 - Phase 2: Shrink - Release the Locks
- Variation of 2PL:
 - To get Strict Schedules (see pp757-759, Chap 21.4.2), Use Strict 2PL
 - Strict 2PL if T commits or aborts before any unlock() operations.
 - Strict 2PL does not avoid deadlocks.

Lock Release

Two-Phase locking: All locks are acquired before any lock is released



Strict: Transaction holds all locks until completion



Correctness of Strict Two-Phase Locking

- Intuition: Active transactions cannot have executed operations that do not commute (since locks required for non-commutative operations conflict)
- Hence, a schedule produced by a two-phase locking concurrency control is serializable since operations of concurrent transactions can always be reordered to produce a serial schedule

Non-Strict Concurrency Controls

- Non-strict controls: locks can be released before completion
- Problem: $w_1(x) \ u_1(x) \ r_2(x) \ w_2(y) \ commit_2 \ abort_1$
 - Although $abort_1$ rolls x back, the new value of y may have been affected
 - $T1$ has an effect even though it is aborted
 - Hence, atomicity is violated

Deadlock and Livelock Problems

- Livelocks or Deadlocks - possible with 2PL
- When a transaction can hold locks and request another lock (*e.g.*, in two-phase locking), a cycle of waiting transactions can result:
 - $r_lock1(Y), r1(Y), r_lock2(X), r2(X), w_lock1(X), w_lock2(Y)$
 - pp 786 Fig 22.5
- A transaction in the cycle must be aborted by DBMS (since transactions will wait forever)
- DBMS uses deadlock detection algorithms or timeout to deal with this problem

Lock based CCT: Solving Deadlock Problems

- Priority based on age (timestamp of first submission)
 - E.g. $TS(T1) < TS(T2)$
- Solution 1: Prevent Deadlocks
 - Conservative 2PL if T acquires all locks before T starts execution.
 - No deadlocks, if all locks requested together (no hold & wait).
- Sol. 2: Detect and Resolve Deadlocks
 - Wait-for graph (nodes = Ts, edge = wait-for dependency)
 - cycle in WFG \Rightarrow deadlock
 - Break deadlock by aborting T_i to break cycles in WFG
- Sol. 3: Timeouts: T_i aborts if all locks are not granted within timeout!
 - wait-die (new time-stamp) or wait-wound (no change in time-stamp)

CCT 2: Timestamp Ordering(pp594 18.2)

- IDEA: Make schedule eqv. to serial schedule
 - defined by the ordering of time-stamps (transX start time) of transactions
 - $TS(T)$: timestamp of transX T
 - No deadlock
- Thus, Abort T_i operate on data-item X, such that
 - T_i reads /writes X with $write_TS(X) > TS(T_i)$
 - OR T_i writes X with $read_TS(X) > TS(T_i)$
 - WHERE
 - ... $TS(T_i)$ = start time of a transaction T_i
 - ... $read_TS(X)$ = largest $TS(T_i)$ for any T_i that has already read X.
 - ... $write_TS(X)$ = largest $TS(T_i)$, any T_i has already written X.
- Timestamp Ordering ensures serializability
 - Note: Potential conflict detected before operation and avoided.
 - Conflict eqv. schedule = order T_i in S by their timestamps.
- Comparison of TSO, 2PL, CGS
 - 2PL/TSO limit concurrency: rules out some serializable schedules
 - 2PL and TSO produce different subsets of CGS

Other CCTs: Versions

- *18.3 Multiversion CCT*
 - Uses view serializability, not conflict serializability
 - $W_i(X)$ writes a new version of X , old version of X remains
 - Some conflicting $r_i(X)$ can read older versions
 - Can be based on timestamps or 2PL
 - Oracle: timestamped-versions for read and 2PL for write

Other CCTs: Optimistic

- *22.4 Optimistic Validation Based CCT*
 - Idea: do all the checking at one time (commit time)
 - ...No side-effect to database before validation
 - Three phases: read + compute, validate, write
 - Abort T_i if conflict with committed / validated Xactions
 - Wins if there is small degree of conflicts across concurrent T_s
 - Too many aborts if too many conflicts