

Algebra

- Based on operators and a domain of values
- Operators map arguments from domain into another domain value
- Hence, an expression involving operators and arguments produces a value in the domain.
- We refer to the expression as a query and the value produced as the result of that query

1

Relational Algebra

- Domain: set of relations
- Basic operators:
 - Select (σ)
 - project (Π)
 - union (\cup)
 - set difference ($-$)
 - Cartesian product (\times)
- Derived operators: set intersection, division, join
- Procedural: Relational expression specifies query by describing an algorithm (the sequence in which operators are applied) for determining the result of an expression

2

Select Operator

- Produce table containing subset of rows of argument table satisfying condition

$$\sigma_{\langle condition \rangle}(\langle relation \rangle)$$

- Example: $\sigma_{Hobby='stamps'}(Person)$

Id	Name	Address	Hobby
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

Person

Id	Name	Address	Hobby
1123	John	123 Main	stamps
9876	Bart	5 Pine St	stamps

3

Selection Condition

- Operators: $<$, \leq , \geq , $>$, $=$, \neq
- Simple selection condition:
 - $\langle attribute \rangle$ operator $\langle constant \rangle$
 - $\langle attribute \rangle$ operator $\langle attribute \rangle$
- $\langle condition \rangle$ AND $\langle condition \rangle$
- $\langle condition \rangle$ OR $\langle condition \rangle$
- NOT $\langle condition \rangle$

4

Selection Condition - Examples

- $\sigma_{Id > 3000 \text{ Or } Hobby = 'hiking'}(Person)$
- $\sigma_{Id > 3000 \text{ AND } Id < 3999}(Person)$
- $\sigma_{NOT(Hobby = 'hiking')}(Person)$
- $\sigma_{Hobby \neq 'hiking'}(Person)$

5

Project Operator

- Produce table containing subset of columns of argument table

$\Pi_{\langle \text{attribute-list} \rangle}(\langle \text{relation} \rangle)$

- Example: $\Pi_{Name, Hobby}(Person)$

Id	Name	Address	Hobby
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

Person

Name	Hobby
John	stamps
John	coins
Mary	hiking
Bart	stamps

6

Project Operator

Example: $\Pi_{Name, Address} (Person)$

Id	Name	Address	Hobby
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

Name	Address
John	123 Main
Mary	7 Lake Dr
Bart	5 Pine St

Person

Result is a table (no duplicates)

7

Expressions

$\Pi_{Id, Name} (\sigma_{Hobby='stamps' \text{ OR } Hobby='coins'} (Person))$

Id	Name	Address	Hobby
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

Person

Id	Name
1123	John
9876	Bart

8

Set Operators

- Relation is a set of tuples \Rightarrow set operations should apply
- Result of combining two relations with a set operator is a relation \Rightarrow all its elements must be tuples having same structure
- Hence, scope of set operations limited to *union compatible relations*

9

Union Compatible Relations

- Two relations are union compatible if
 - Both have same number of columns
 - Domains of the corresponding attributes in both relations must be compatible
- Union compatible relations can be combined using union, intersection, and set difference

10

Example

Tables:

Person (SSN, Name, Address, Hobby)

Professor (Id, Name, Office, Phone)

are not union compatible. However

$\Pi_{Name}(Person)$ and $\Pi_{Name}(Professor)$

are union compatible and

$\Pi_{Name}(Person) - \Pi_{Name}(Professor)$

makes sense.

11

Cartesian Product

- If R and S are two relations, $R \times S$ is the set of all concatenated tuples $\langle x, y \rangle$, where x is a tuple in R and y is a tuple in S
 - (R and S need not be union compatible)
- $R \times S$ is expensive to compute:
 - Factor of two in the size of each row
 - Quadratic in the number of rows

a	b
x1	x2
x3	x4

R

c	d
y1	y2
y3	y4

S

a	b	c	d
x1	x2	y1	y2
x1	x2	y3	y4
x3	x4	y1	y2
x3	x4	y3	y4

$R \times S$

12

Renaming

- Result of expression evaluation is a relation
- Attributes of relation must have distinct names. This is not guaranteed with Cartesian product
 - e.g., suppose in previous example $a = c$
- Renaming operator tidies this up. To assign the names A_1, A_2, \dots, A_n to the attributes of the n column relation produced by *expression* use *expression* $[A_1, A_2, \dots, A_n]$

13

Example

Transcript (StudId, CrsCode, Semester, Grade)

Teaching (ProfId, CrsCode, Semester)

$\Pi_{StudId, CrsCode} (Transcript) [StudId, SCrsCode] \times$
 $\Pi_{ProfId, CrsCode} (Teaching) [ProfId, PCrscode]$

This is a relation with 4 attributes:

StudId, SCrsCode, ProfId, PCrsCode

14

Derived Operation: Join

The expression :

$$\sigma_{join-condition'}(R \times S)$$

where *join-condition'* is a *conjunction* of terms:

$$A_i \text{ oper } B_i$$

in which A_i is an attribute of R , B_i is an attribute of S , and *oper* is one of $=, <, >, \geq, \neq, \leq$, is referred to as the (theta) join of R and S and denoted $R \bowtie S$.

$$R \bowtie_{join-condition} S$$

Where *join-condition* and *join-condition'* are (roughly) the same ... 15

Join and Renaming

- **Problem:** R and S might have attributes with the same name – in which case the Cartesian product is not defined
- **Solution:**
 - Rename attributes prior to forming the product and use new names in *join-condition'*.
 - Common attribute names are qualified with relation names in the result of the join

Theta Join – Example

Output the names of all employees that earn more than their managers.

$$\Pi_{Employee.Name} (Employee \bowtie_{MngrId=Id \text{ AND } Salary > Salary} Manager)$$

The join yields a table with attributes:

Employee.Name, Employee.Id, Employee.Salary, MngrId
Manager.Name, Manager.Id, Manager.Salary

17

Equijoin Join - Example

Equijoin: Join condition is a conjunction of *equalities*.

Student				Transcript			
Id	Name	Addr	Status	StudId	CrsCode	Sem	Grade
111	John	111	CSE305	S00	B
222	Mary	222	CSE306	S99	A
333	Bill	333	CSE304	F99	A
444	Joe				

Mary	CSE306
Bill	CSE304

The equijoin is commonly used since it combines related data in different relations.

18

Equijoin Join - Example

Equijoin: Join condition is a conjunction of *equalities*.

$$\Pi_{Name, CrsCode}(Student \bowtie_{Id=StudId \text{ and } Grade='A'} Transcript)$$

Student			
Id	Name	Addr	Status
111	John
222	Mary
333	Bill
444	Joe

Transcript			
StudId	CrsCode	Sem	Grade
111	CSE305	S00	B
222	CSE306	S99	A
333	CSE304	F99	A

Mary	CSE306
Bill	CSE304

The equijoin is commonly used since it combines related data in different relations.

19

Climbers (C2):

CId	CName	Skill	Age
123	Edmund	EXP	80
214	Arnold	BEG	25
313	Bridget	EXP	33
212	James	MED	27

Climbs (C1):

CId	RId	Date	Duration
123	1	10/10/88	5
123	3	11/08/87	1
313	1	12/08/89	5
214	2	08/07/92	2
313	1	06/07/94	3

$$\sigma_{CId:1=CId:2}(Climbs \times Climbers):$$

C1.CId	RId	Date	Duration	C2.CId	CName	Skill	Age
123	1	10/10/88	5	123	Edmund	EXP	80
123	3	11/08/87	1	123	Edmund	EXP	80
313	1	12/08/89	5	313	Bridget	EXP	33
214	2	08/07/92	2	214	Arnold	BEG	25
313	1	06/07/94	3	313	Bridget	EXP	33

20

Natural JOIN (*)

- In an EQUIJOIN $R \leftarrow R_1 \bowtie_c R_2$, the join attribute of R_2 appear *redundantly* in the result relation R.
- In a NATURAL JOIN, the *redundant join attributes* of R_2 are *eliminated* from R. The equality condition is *implied* and need not be specified.

Climbs \bowtie Climbers :

CId	RId	Date	Duration	CName	Skill	Age
123	1	10/10/88	5	Edmund	EXP	80
123	3	11/08/87	1	Edmund	EXP	80
313	1	12/08/89	5	Bridget	EXP	33
214	2	08/07/92	2	Arnold	BEG	25
313	1	06/07/94	3	Bridget	EXP	33

21

Natural Join

- Special case of equijoin:
 - join condition equates *all* and *only* those attributes with the same name (condition doesn't have to be explicitly stated)
 - duplicate columns eliminated from the result

Transcript (*StudId*, *CrsCode*, *Sem*, *Grade*)

Teaching (*ProfId*, *CrsCode*, *Sem*)

$$\begin{aligned}
 & \text{Transcript} \bowtie \text{Teaching} = \\
 & \pi_{\text{StudId}, \text{Transcript.CrsCode}, \text{Transcript.Sem}, \text{Grade}, \text{ProfId}} (\\
 & \quad \text{Transcript} \bowtie_{\text{CrsCode}=\text{CrsCode} \text{ AND } \text{Sem}=\text{Sem}} \text{Teaching}) \\
 & \quad [\text{StudId}, \text{CrsCode}, \text{Sem}, \text{Grade}, \text{ProfId}]_{22}
 \end{aligned}$$

Natural Join (con't)

- More generally:

$$R \bowtie S = \pi_{attr-list} (\sigma_{join-cond} (R \times S))$$

where

$$attr-list = attributes(R) \cup attributes(S)$$

(duplicates are eliminated) and *join-cond* has the form:

$$A_1 = A_1 \text{ AND } \dots \text{ AND } A_n = A_n$$

where

$$\{A_1 \dots A_n\} = attributes(R) \cap attributes(S)$$

23

Natural Join Example

- List all Id's of students who took at least two different courses:

Natural Join Example

- List all Id's of students who took at least two different courses:

$$\Pi_{StudId} (\sigma_{CrsCode \neq CrsCode2} ($$

$$Transcript * Transcript [StudId, CrsCode2, Sem2, Grade2]))$$

(don't join on CrsCode, Sem,
and Grade attributes)


25

Division

- Goal: Produce the tuples in one relation, r , that match *all* tuples in another relation, s
 - $r (A_1, \dots A_n, B_1, \dots B_m)$
 - $s (B_1 \dots B_m)$
 - r/s , with attributes $A_1, \dots A_n$, is the set of all tuples $\langle a \rangle$ such that for every tuple $\langle b \rangle$ in s , $\langle a, b \rangle$ is in r
- Can be expressed in terms of projection, set difference, and cross-product

26

Division (con't)



27

Division - Example

- List the Ids of students who have passed *all* courses that were taught in spring 2000

28

Division - Example

- List the Ids of students who have passed *all* courses that were taught in spring 2000
- *Numerator*: StudId and CrsCode for every course passed by every student
 - $\pi_{StudId, CrsCode}(\sigma_{Grade \neq 'F'}(Transcript))$
- *Denominator*: CrsCode of all courses taught in spring 2000
 - $\Pi_{CrsCode}(\sigma_{Semester='S2000'}(Teaching))$
- Result is *numerator/denominator*

29

Complete Set of Relational Algebra Operations

- the set $\{\sigma, \Pi, U, -, \bowtie\}$ is called a *complete set* of relational algebra operations. Any query language *equivalent to* these operations is called **relationally complete**.
- All the basic operations discussed so far can be described as a sequence of *only* the above set
- Additional operations were not part of the *original* relational algebra
 - Aggregate functions (SUM, AVG), grouping
 - Outer Join, outer union

30

Additional Relational Operations

- Aggregate Functions and Grouping
- Outer JOIN and Outer UNION

Aggregate functions

- Functions are often applied to sets of values or sets of tuples in DB applications
 - SUM, COUNT, AVERAGE, MIN, MAX
 - $\langle \text{grouping attributes} \rangle \mathbf{F} \langle \text{function list} \rangle (\mathbf{R})$
 - Grouping attributes are optional

- **Example:**

List the average salary of **all** employees (no grouping needed):

For each department, retrieve the department number, the number of employees, and the average salary:

33

Aggregate functions

- Functions are often applied to sets of values or sets of tuples in DB applications
 - SUM, COUNT, AVERAGE, MIN, MAX
 - $\langle \text{grouping attributes} \rangle \mathbf{F} \langle \text{function list} \rangle (\mathbf{R})$
 - Grouping attributes are optional

- **Example:**

List the average salary of **all** employees (no grouping needed):

$\mathbf{R}(\text{AVGSAL}) \leftarrow \mathbf{F}_{\text{AVG SALARY}} (\mathbf{EMPLOYEE})$

For each department, retrieve the department number, the number of employees, and the average salary:

$\mathbf{R}(\text{DNO}, \text{NUMEMPS}, \text{AVGSAL}) \leftarrow$

$\mathbf{DNO} \mathbf{F}_{\text{COUNT SSN, AVERAGE SALARY}} (\mathbf{EMPLOYEE})$

- DNO is called the **grouping attribute**

34

Figure 7.16 An illustration of the AGGREGATE FUNCTION operation. (a) $R(DNO, NO_OF_EMPLOYEES, AVERAGE_SAL) \leftarrow \pi_{DNO} \sigma_{COUNT_SSN, AVERAGE_SALARY}(EMPLOYEE)$. (b) $\pi_{DNO} \sigma_{COUNT_SSN, AVERAGE_SALARY}(EMPLOYEE)$. (c) $\sigma_{COUNT_SSN, AVERAGE_SALARY}(EMPLOYEE)$.

(a)

	DNO	NO_OF_EMPLOYEES	AVERAGE_SAL
	5	4	33250
	4	3	31000
	1	1	55000

(b)

DNO	COUNT_SSN	AVERAGE_SALARY
5	4	33250
4	3	31000
1	1	55000

(c)

COUNT_SSN	AVERAGE_SALARY
8	35125

© Addison Wesley Longman, Inc. 2000, Elmasri/Navathe, Fundamentals of Database Systems, Third Edition

35

Outer join

- In a regular EQUIJOIN or NATURAL JOIN operation, tuples in R1 or R2 that do not have matching tuples in the other relation *do not appear in the result*
- Tuples with null in the join attributes are also eliminated.
- Some queries require all tuples in R1 (or R2 or both) to appear in the result
- When no matching tuples are found, **nulls** are placed for the missing attributes

36

Outer join (Cont.)

- **LEFT OUTER JOIN:**
 - $R1 \text{ left} \bowtie R2$ lets every tuple in $R1$ appear in the result
- **RIGHT OUTER JOIN:**
 - $R1 \bowtie \text{right} R2$ lets every tuple in $R2$ appear in the result
- **FULL OUTER JOIN:**
 - $R1 \text{ left} \bowtie \text{right} R2$ lets every tuple in $R1$ or $R2$ appear in the result

37

Figure 7.18 The LEFT OUTER JOIN operation.

RESULT	FNAME	MINIT	LNAME	DNAME
	John	B	Smith	null
	Franklin	T	Wong	Research
	Alicia	J	Zelaya	null
	Jennifer	S	Wallace	Administration
	Ramesh	K	Narayan	null
	Joyce	A	English	null
	Ahmad	V	Jabbar	null
	James	E	Borg	Headquarters

38

Outer UNION

- If two relations are **not union compatible**
- Partial compatible
- example